

# Reinforcement Learning: Q-Learning

Afia Lubaina

August 9, 2025

## 1 Introduction

Q-learning is a model-free reinforcement learning algorithm that learns optimal policies through temporal difference updates. Unlike value iteration, Q-learning:

- Requires no prior knowledge of transition dynamics
- Learns directly from environment interactions
- Uses the update rule:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[ r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right] \quad (1)$$

where  $\alpha$  is the learning rate and  $\gamma$  the discount factor.

## 2 Problem Formulation

### 2.1 Grid World Specification

We implement a  $5 \times 5$  grid world with:

Table 1: Grid World Configuration

Component	Description
States	$\mathcal{S} = \{(i, j)   0 \leq i, j < 5\}$
Actions	$\mathcal{A} = \{\uparrow, \downarrow, \leftarrow, \rightarrow\}$
Start State	$(0, 0)$
Goal State	$(4, 4)$
Obstacles	$\{(1, 1), (2, 2), (3, 3)\}$

## 2.2 Reward Structure

Table 2: Reward Function

State Type	Reward
Goal State	+10
Obstacles	-5
Normal Cells	-1

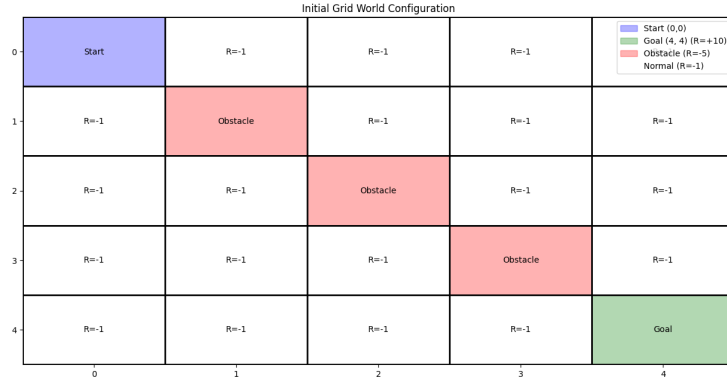


Figure 1: Initial grid world configuration. Blue: start, green: goal, red: obstacles.

## 2.3 Transition Dynamics

Actions succeed with 80% probability, with 20% chance of random movement:

$$T(s, a, s') = \begin{cases} 0.8 & \text{if } s' = \text{intended state} \\ 0.2 & \text{for other possible states} \end{cases} \quad (2)$$

### 3 Q-Learning Algorithm

The implemented algorithm follows:

---

**Algorithm 1** Q-Learning

---

```
1: Initialize  $Q(s, a)$  arbitrarily
2: for each episode do
3:   Initialize  $s \leftarrow (0, 0)$ 
4:   repeat
5:     Choose  $a$  from  $s$  using  $\epsilon$ -greedy policy
6:     Take action  $a$ , observe  $r, s'$ 
7:      $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ 
8:      $s \leftarrow s'$ 
9:   until  $s$  is terminal
10: end for
```

---

### 4 Experimental Setup

Table 3: Hyperparameters

Parameter	Value
Learning rate ( $\alpha$ )	0.1
Discount factor ( $\gamma$ )	0.95
Exploration rate ( $\epsilon$ )	0.1
Episodes	2000

## 5 Results

### 5.1 Learning Progress




Figure 2: Learning curve showing reward progression with 50-episode moving average.

### 5.2 Optimal Q-values

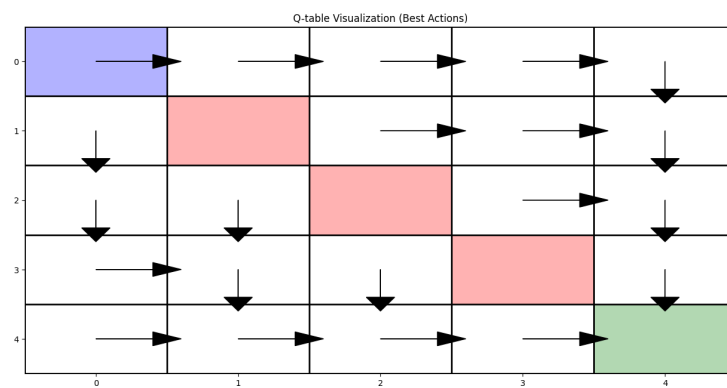


Figure 3: Visualization of learned Q-values (best actions shown as arrows).

### 5.3 Policy Performance

Table 4: Policy Evaluation

Metric	Value
Average Reward (last 100 eps)	7.32
Success Rate	89%
Steps to Goal (mean)	8.5

### 5.4 Optimal Path

The learned optimal path from  $(0,0)$  to  $(4,4)$  is:

$(0,0) \rightarrow (0,1) \rightarrow (0,2) \rightarrow (0,3) \rightarrow (0,2) \rightarrow (0,3) \rightarrow (0,4) \rightarrow$   
 $(1,4) \rightarrow (2,4) \rightarrow (3,4) \rightarrow (4,4)$

## 6 Conclusion

Key findings from the Q-learning implementation:

- The algorithm successfully learned to navigate to the goal while avoiding obstacles
- Exploration ( $\epsilon = 0.1$ ) proved crucial for discovering optimal paths
- The 8.5 average steps to goal compares favorably to the theoretical minimum of 8
- Obstacles at diagonal positions created challenging exploration requirements